

Homework 4

Ido Nutov & Tuvy Lemberg

July 28, 2025

1

We study two realizations of the action of the symmetric group S_n for finite n :

- **Index action:** S_n acts on $X = \{1, 2, \dots, n\}$ by $g \cdot i = g(i)$.
- **Coordinate action:** S_n acts on $v \in \mathbb{R}^n$ by permuting coordinates: $(gv)_i = v_{g^{-1}(i)}$.

These actions are equivalent, as the coordinate action is the index action applied to the vector indices. The standard permutation representation is $\rho : S_n \rightarrow GL(n, \mathbb{R})$, with $\rho(g)_{i,j} = \delta_{j,g(i)}$.

Let $G \leq S_n$. Let's look at the G -equivariant linear maps $T : \mathbb{R}^n \rightarrow \mathbb{R}^n$, such that

$$T\rho(g)v = \rho(g)Tv \quad (1)$$

for all $g \in G, v \in \mathbb{R}^n$.

We can solve (1) in two ways:

- **Orbits of Indexes:** Looking on the condition (1) applied to the indexes of T .
- **Character Theory:** Looking on the condition (1) applied to the representation ρ of G .

Orbits of Indexes Since G acts on the indexes of T , we have

$$\rho(g)T\rho(g)^{-1} = T, \quad (2)$$

for all $g \in G$ lead to the condition

$$T_{i,j} = T_{g^{-1}(i),g^{-1}(j)}. \quad (3)$$

This holds for all $g \in G$ and $i, j \in X$, therefore, for each pair of indexes (i, j) , the value $T_{i,j}$ is constant on the orbit $Orb_G((i, j))$. Leading to the conclusion that the number of free parameters in T (the dimensionality $\dim HOM_G(\mathbb{R}^n, \mathbb{R}^n)$) is equal to the number of orbits of the action of G on X^2 . The number of orbits can be computed using Burnside's lemma as in Tutorial 12 Problem 2

$$|X^2/G| = \frac{1}{|G|} \sum_{g \in G} |X^g|^2 \quad (4)$$

where $X^g = \{x \in X \mid g \cdot x = x\}$ is the set of fixed elements of g .

Character Theory In the course we have seen the proof that the dimensionality of the space of G -linear transformations $HOM_G(\mathbb{R}^n, \mathbb{R}^n)$ is equal to

$$\dim HOM_G(\mathbb{R}^n, \mathbb{R}^n) = \frac{1}{|G|} \sum_{g \in G} \text{tr}(\rho(g))^2. \quad (5)$$

From condition (2) we have

$$(\rho(g) \otimes \rho(g))\text{vec}(T) = \text{vec}(T). \quad (6)$$

We can see that $\text{vec}(T)$ is an eigenvector of $\rho(g) \otimes \rho(g)$ with eigenvalue 1 for all $g \in G$. Moreover, let $\{O_1, \dots, O_k\}$ be the orbits of the action of G on X^2 , where k is the number of orbits; and let $\mathbf{1}_{O_i} \in \mathbb{R}^{n \times n}$ be the indicator matrix of the orbit O_i , i.e. $\mathbf{1}_{O_i}(i, j) = 1$ if $(i, j) \in O_i$ and 0 otherwise. Then, $\text{vec}(\mathbf{1}_{O_i})$ must be an eigenvector of $\rho(g) \otimes \rho(g)$ with eigenvalue 1 for all $g \in G$ as well, since it is constant when g acts on it. Define $U^G = \{v \in \mathbb{R}^{n \times n} \mid (\rho(g) \otimes \rho(g))v = v, \text{ for all } g \in G\}$, as the fixed subspace of $\mathbb{R}^{n \times n}$ under the action of G . U^G is isomorphic to the space of G -linear transformations $HOM_G(\mathbb{R}^n, \mathbb{R}^n)$. In the course we saw that the dimension of U^G is equal $\text{tr}(\pi_{U^G})$, where π_{U^G} is the projection operator onto U^G . We showed that $\pi_{U^G} = \frac{1}{|G|} \sum_{g \in G} \rho(g) \otimes \rho(g)$, and therefore we have

$$\dim HOM_G(\mathbb{R}^n, \mathbb{R}^n) = \frac{1}{|G|} \sum_{g \in G} \text{tr}(\rho(g) \otimes \rho(g)) = \frac{1}{|G|} \sum_{g \in G} \text{tr}(\rho(g))^2. \quad (7)$$

This is the same as (5), and each eigenvalue of $\rho(g) \otimes \rho(g)$ corresponds to an orbit of the action of G on X^2 . It is clear that each orbit corresponds to a unique eigenvector of $\rho(g) \otimes \rho(g)$, but the other way need to be explained. If we have an eigenvector v of $\rho(g) \otimes \rho(g)$, it must satisfy (3), which means that it must be constant on the orbits, therefore it can be represented as a linear combination of the $\{\text{vec}(O_i)\}_{i=1}^k$, and therefore it cannot add a dimension to the eigenvalues corresponding to the orbits. In the case of index action X^g corresponding to the fixed points of g acting on the indexes. In the case of coordinate action, U^G corresponds to the fixed vector space of G acting on the vectorization of the matrices T .

2

2.1

The natural group here is the cross product of the permutation groups $S_n \times S_m$ acts on the Netflix rating matrices $A \in \mathbb{R}^{n \times m}$. It can be explained by the fact that the order of the m movies and the order of the n users is arbitrary and both can be changed independently.

2.2

Since we look at the G -linear transformation from matrices in $\mathbb{R}^{n \times m}$ to itself, it will be more convenient to consider the vectorization of these matrices. Thus $\text{vec}(A) = [a_{1,1}, \dots, a_{n,1}, \dots, a_{1,m}, \dots, a_{n,m}]^T \in \mathbb{R}^{nm}$. Let $g = (g_l, g_r) \in S_n \times S_m$ be a group element where $g_l \in S_n$ is a left action that acts on the rows of A and $g_r \in S_m$ acts on the columns of A as follows:

$$gA = g_l A g_r^T = P_l A P_r^T, \quad (8)$$

where $P_l \in \mathbb{R}^{n \times n}$ and $P_r \in \mathbb{R}^{m \times m}$ are the permutations of g_l and g_r , respectively. Define $v = \text{vec}(A) \in \mathbb{R}^{nm}$. We are actually looking on the isomorphic group $G \leq S_{nm}$, which acts on the indexes of the vectorization of A . We will use G and $S_n \times S_m$ interchangeably. In index notation we have $(gA)_{i,j} = A_{g_l^{-1}(i), g_r^{-1}(j)}$. We are looking at the affine transformation that is equivariant, thus $f(v) = Wv + b$, that fulfills for all $g \in S_n \times S_m$ and $v \in \mathbb{R}^{nm}$

$$\begin{aligned} f(gv) &= W(gv) + b = g(Wv + b) \\ &\Rightarrow Wgv + b = gWv + gb \end{aligned} \quad (9)$$

Since (9) needs to be fulfilled for all $v \in \mathbb{R}^{nm}$ it leads to $Wg = gW$ and equivalently $gWg^T = W$. Moreover, it requires that $gb = b$ for all $g \in G$. We have

$$gA = g_l A g_r^T = P_l A P_r^T. \quad (10)$$

By vectorizing the above we obtain

$$\text{vec}(P_l A P_r^T) = (P_r \otimes P_l) \text{vec}(A) \Rightarrow P_{nm} = (P_r \otimes P_l) \quad (11)$$

For b , we can see from (9) that this means that b is constant, $b = c \cdot \mathbf{1}$. This can be seen by applying the inverse vectorization operation on b . Since the condition (9) requires that b will be invariant to the action of g , and by acting on b we can replace each index $\text{vec}^{-1}(b_{i,j})$ with any other index. Therefore, all values of b required to be the same.

Let us look at the action of the permutation $g^{-1} = (g_l^{-1}, g_r^{-1})$ on the vector v . We take the inverse g^{-1} for convenience of notation. Originally, the element with the index $A_{i,j}$ moves to $v_{i+n(j-1)}$, therefore

$$(g^{-1}A)_{i,j} = A_{g_l(i), g_r(j)} = v_{g_l(i) + n(g_r(j)-1)}. \quad (12)$$

Consequently, for an index $1 \leq k \leq nm$, define $j = k/n$ as the truncating integer division and $i = k \% n$ as the modulo operation, thus $(g^{-1}v)_k = v_{g_l(i) + n(g_r(j)-1)}$. Let k, l be indexes in $[1, \dots, nm]$ with the decomposition $k = i_1 + n(j_1 - 1)$ and $l = i_2 + n(j_2 - 1)$. Then $(g^{-1}Wg)_{k,l} = W_{g_l(i_1) + n(g_r(j_1)-1), g_l(i_2) + n(g_r(j_2)-1)}$.

2.2.1 Orbit analysis

Let us look at the possible orbits in the index set $X^2 = \{1, 2, \dots, nm\}^2$. Indexes (k', l') can be on the same orbits as (k, l) if and only if there exists a permutation $(g_l, g_r) \in S_n \times S_m$ such that

$$\begin{cases} k' &= g_l(i_1) + n(g_r(j_1) - 1) \\ l' &= g_l(i_2) + n(g_r(j_2) - 1) \end{cases} \quad (13)$$

Define i_3, i_4, j_3, j_4 such that $k' = i_3 + n(j_3 - 1)$ and $l' = i_4 + n(j_4 - 1)$. It is clear that we can create equivalence groups by the rule:

$$(i_1 \neq i_2) \wedge (j_1 \neq j_2). \quad (14)$$

It is clear that $i_3 = i_4$ holds if and only if $i_1 = i_2$, since by applying g_l we have $g_l(i_1) = g_l(i_2) = i_3 = i_4$, and since g_l is a bijection. The same goes for $j_3 = j_4$.


These are the orbits and the explanation for why they are the only orbits is as follows: if $i_1 = i_2$ then we can move to any other k', l' that has $i_3 = i_4$ by applying the right $g_l(i_1) = i_3$. If $i_1 \neq i_2$

we can move to any other indexes $i_3 \neq i_4$ since S_n is 2-transitive. The same goes for the indexes j . Therefore, we will have only these 4 orbits:

$$\begin{aligned} O_{=i=j} &: i_1 = i_2, j_1 = j_2 \\ O_{=i \neq j} &: i_1 = i_2, j_1 \neq j_2 \\ O_{\neq i=j} &: i_1 \neq i_2, j_1 = j_2 \\ O_{\neq i \neq j} &: i_1 \neq i_2, j_1 \neq j_2 \end{aligned}$$

2.3

For $n = m = 3$ we will obtain $W \in \mathbb{R}^{9 \times 9}$ and $b \in \mathbb{R}^9 = c \cdot \mathbf{1}$



2.4

In the case of a linear invariant layer, for all $v \in \mathbb{R}^{nm}$ and $g \in G$ we have:

$$\begin{aligned} f(gv) &= f(v) \\ \Rightarrow w^T gv + b &= w^T v + b \\ \Rightarrow g^T w &= w \end{aligned}$$

For $w \in \mathbb{R}^{nm}$, this condition is equivalent to the condition on b in the previous clause. Therefore, $w = c_2 \cdot \mathbf{1}$.

2.5

In the case of 3-tensors with the symmetric group is $S_{N_1} \times S_{N_2} \times S_{N_3}$ applied independently to each index axis of the tensor. We have:

$$v = \text{vec}(A) \quad v_k = a_{i_1, i_2, i_3}, \quad (15)$$

where $k = i_1 + N_1(i_2 - 1) + N_1N_2(i_3 - 1)$. Define $k = i_1^k + N_1(i_2^k - 1) + N_1N_2(i_3^k - 1)$ and $l = i_1^l + N_1(i_2^l - 1) + N_1N_2(i_3^l - 1)$. We will have the same derivation as in the previous clause 2.2, thus

$$W_{k,l} = W_{g(k),g(l)}$$

The indexes (k, l) can transform to

$$\begin{aligned} k' &= g_1(i_1^k) + N_1(g_2(i_2^k) - 1) + N_1N_2(g_3(i_3^k) - 1) \\ l' &= g_1(i_1^l) + N_1(g_2(i_2^l) - 1) + N_1N_2(g_3(i_3^l) - 1) \end{aligned}$$


where the equivalence relations follow the same pattern as before, thus the orbits are

$$\bigwedge_{n=1}^3 i_n^k \neq i_n^l. \quad (16)$$

As a result, the orbits are:

$$\begin{aligned} O_{=i_1=i_2=i_3} &: i_1^k = i_1^l, i_2^k = i_2^l, i_3^k = i_3^l \\ O_{=i_1=i_2 \neq i_3} &: i_1^k = i_1^l, i_2^k = i_2^l, i_3^k \neq i_3^l \\ O_{=i_1 \neq i_2=i_3} &: i_1^k = i_1^l, i_2^k \neq i_2^l, i_3^k = i_3^l \\ &\vdots \\ O_{\neq i_1 \neq i_2 \neq i_3} &: i_1^k \neq i_1^l, i_2^k \neq i_2^l, i_3^k \neq i_3^l \end{aligned}$$

For $N_1 = N_2 = N_3 = 3$, we obtain $W \in \mathbb{R}^{27 \times 27}$ and $b \in \mathbb{R}^{27}$.



[11pt]article
amsmath, amssymb booktabs geometry enumitem pifont margin=2.7cm

3 Code overview

`models.py`

Implements the five architectures required in Q 3: **Canonisation_Net**, **Symmetrization_Net**, **Sampled_Symmetrization_Net**, **Linear_eq_Net** (built from two equivariant layers), and **AugmentedInvariantNet**. Each class exposes `forward(X)` with $X \in \mathbb{R}^{N \times d}$.

`utils.py` • Permutation helpers: `create_permutations_sampled`.

- *Unit tests* for invariance/equivariance (`test_canonization_net, ...`).

- Training routine `train_variance_net` that learns variance purely from permutation augmentations.

`main.py`

Orchestrates the workflow:

- (a) runs each invariance test 50 times and prints the failure rate;
- (b) trains the variance model for 200 epochs on a fixed set of size $N = 50$, $d = 5$;
- (c) retests the trained model for invariance.

Dependencies are limited to `torch` and `numpy`; the script detects a GPU via `torch.cuda.is_available()` and moves tensors accordingly.

4 Permutation invariance sanity check

test Permutation invariance

For every network we apply 50 random permutations $P \in S_N$ to a fixed set X and record a *failure* when the ℓ_∞ distance between the original and permuted outputs exceeds a model-specific tolerance:

- (a,b,d) fail if $\|f(X) - f(PX)\|_\infty > 10^{-5}$;
- (c) fail if $\|f(X) - f(PX)\|_\infty > 2 \times 10^{-1}$
(higher threshold to reflect the Monte Carlo approximation of S_N);
- (e) fail if $\|f(X) - f(PX)\|_\infty > 5 \times 10^{-2}$
(relaxed tolerance to probe the invariance learned via augmentation).

Finally, for each model we report the *failure rate*

$$\text{Failure rate} = \frac{\# \text{ failed permutations}}{\# \text{ permutations tested}},$$

i.e. the fraction of the 50 sampled permutations that violate their respective tolerance.

Network	% noninvariant ↓	Pass?
Canonisation	0.00	✓
Full symmetrisation	0.00	✓
Sampled symmetrisation ($K = 256$)	0.94	✗
S_N -equivariant linear layers	0.00	✓
Learned-invariant (augmentations)	0.28	✓*

Table 1: Failure rate for each network.

Discussion

- **Exact symmetry methods.** Canonisation, full symmetrisation, and linear S_N -equivariant layers are provably invariant; failures are zero within numerical precision.
- **Monte Carlo method (c).** Sampled symmetrisation averages over $K = 0.05 \cdot 8!$ permutations $\approx 0.05 |S_N|$ for $N = 8$ so estimator variance might explain the $\approx 1\%$ failure rate. Increasing K should decrease the error.
- **Learned invariance (e).** The unconstrained MLP acquires symmetry through augmentation; after ~ 250 epochs with ~ 500 augmentations per epoch the failure rate falls below 0.3% .

* Both stochastic models (c and e) pass with a looser tolerance, We intentionally kept a looser threshold here to reveal the *trend* toward invariance and to assess given our time and computation resources how far each model has progressed toward full invariance.